

Project III: Parallel Processing and Multi-Threading

Student: William Duran

Panther ID: 6320114

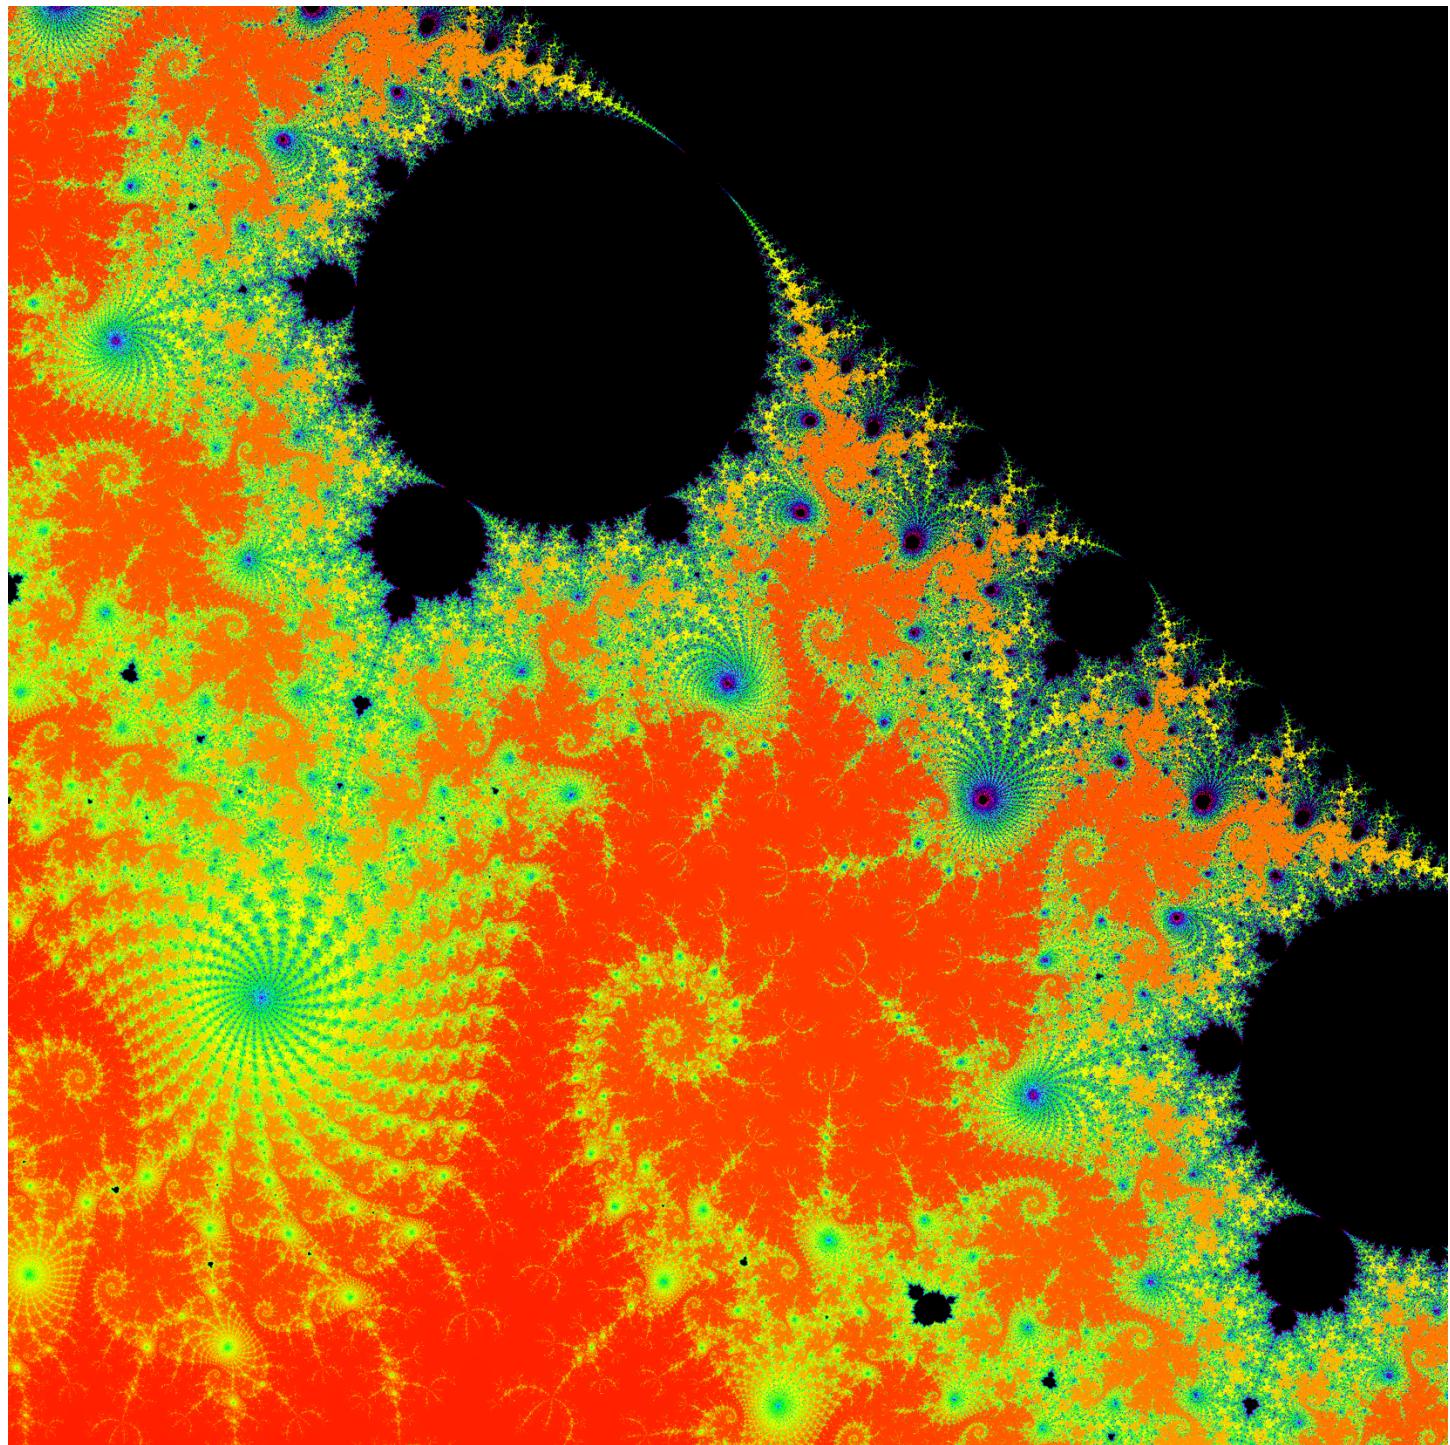


Figure 1 → ./mandel -x -0.5397949000000 -y -0.6095890009734 -s 0.00005 -m 15000 -n 250 -W 2440 -H 2440

Evaluation

- In your own words, briefly explain the purpose of the experiments and the experimental setup. Be sure to clearly state on which machine you ran the experiments, and exactly what your command line arguments were, so that we can reproduce your work in case of any confusion.
- In this project, we are delving into the realms of parallelism using processes and threads to optimize the performance while generating Mandelbrot set images. The setup provided serves as a playground to experiment with both coarse-grained and fine-grained parallelism by modifying the ‘mandel.c’ code to implement multithreading. By measuring the execution time across different thread counts, we aim to evaluate how parallel processing can significantly cut down computation time in such a CPU-intensive task. All the experiments were performed in a 2021 “14” M1 MacBook Pro.

The image I found interesting is mandel -x -0.5397949000000 -y -0.6095890009734 -s 0.0001

The following is the result of executing → ./mandel -x -0.5397949000000 -y -0.6095890009734 -s 0.0001 -m 3500

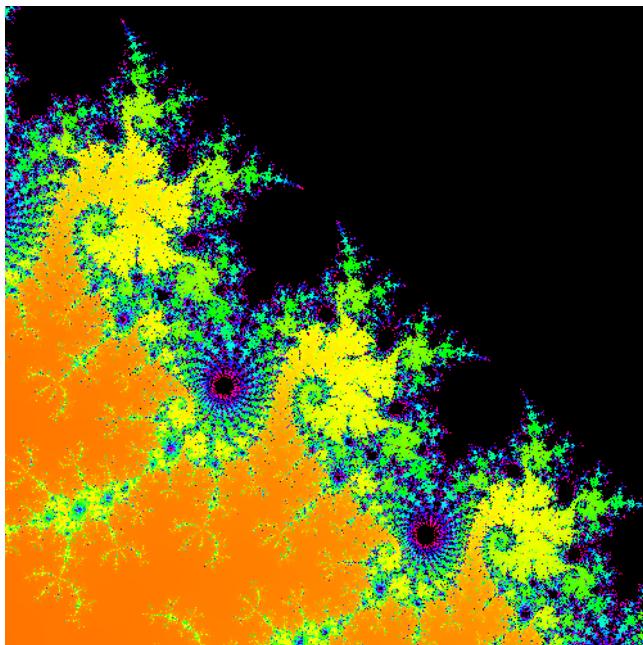
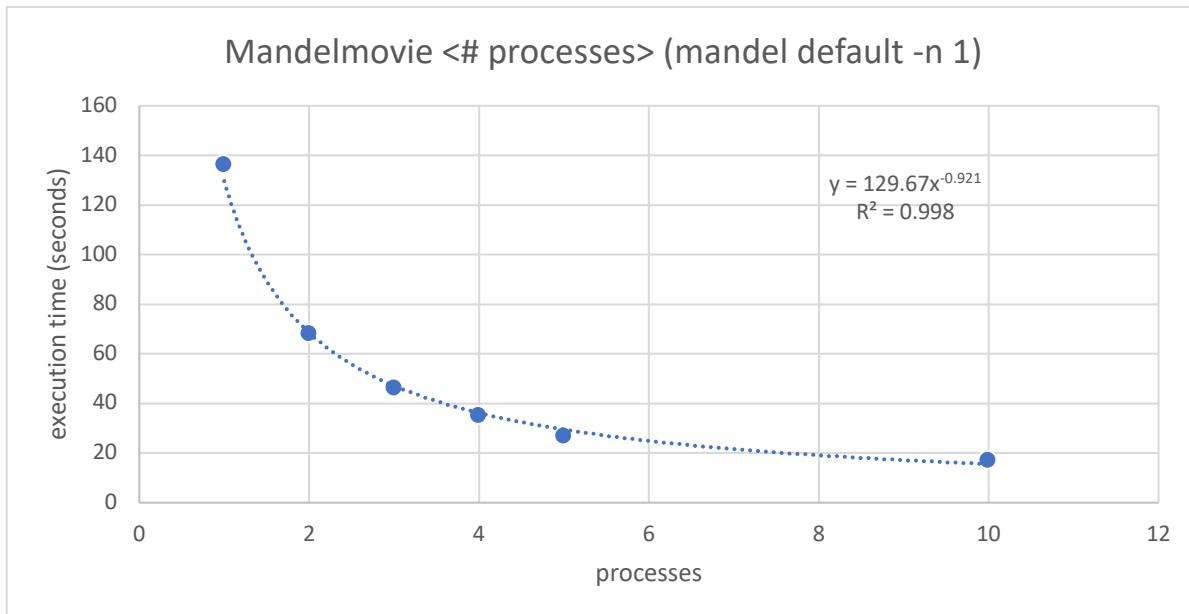


Figure 2 → mandel: x=-0.539795 y=-0.609589 scale=0.000100 max=3500 outfile=mandel.bmp threads=1 Time taken: 5.382914 seconds

- Measure and graph the execution time of mandelmovie for each of 1, 2, 3, 4, 5, and 10 processes running simultaneously. Because each of these will be fairly long running, it will be sufficient to measure each configuration only once. Make sure that you write all of your output files somewhere in /tmp to avoid any timing problems with the network file system.

Mandelmovie <# processes> (mandel default -n 1)	
processes	execution time (seconds)
1	136.198664
2	68.205257
3	46.336647
4	35.23962
5	26.971056
10	16.978963

Table 1 → Mandelmovie <# processes> (mandel default -n 1)



- Explain the shape of the curve. What is the optimal number of processes? Why? Is it possible to have too many processes? Why?
 - Curve Analysis:
 - The graph depicts an exponential decay in execution time with an increase in processes. The most substantial reduction was evident in the initial addition of processes. However, the benefits lessened as more processes were added.
- Optimal Processes:
 - While the exact optimal number is contingent upon specific goals, it's evident that beyond 5 processes, the reduction in execution time is not as pronounced. The jump from 1 to 2 processes yielded a significant decrease, whereas the shift from 5 to 10 was notably less impactful.
- Limitations with Increasing Processes:
 - Introducing too many processes can lead to inefficiencies due to:
 - Overhead in process management.
 - Competition for shared resources.
 - Inherent sequential components in the program, as dictated by Amdahl's Law.
- Conclusion:
 - Although parallelization via multiple processes can enhance performance, there's a threshold beyond which the advantages decline. It's crucial to strike a balance to achieve optimal efficiency.

- For the following two configurations, measure and graph the execution time of multithreaded mandel using 1, 2, 3, 4, 5, 10, 50, and 100 threads. The execution time of these experiments may be upset by other things going on in the machine. So, repeat each measurement five times, and use the fastest time achieved.

- A: `mandel -x -.5 -y .5 -s 1 -m 2000`

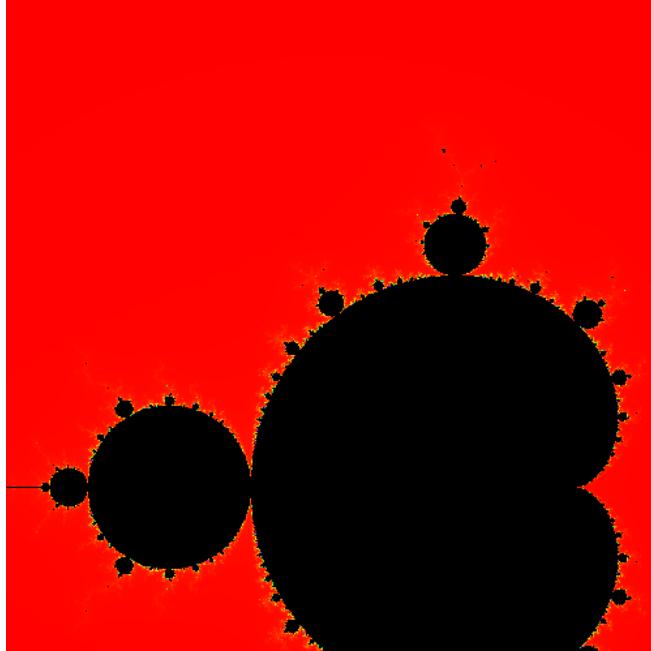


Figure 9 → `./mandel -x -.5 -y .5 -s 1 -m 2000`

- B: `mandel -x 0.2869325 -y 0.0142905 -s .000001 -W 1024 -H 1024 -m 1000`

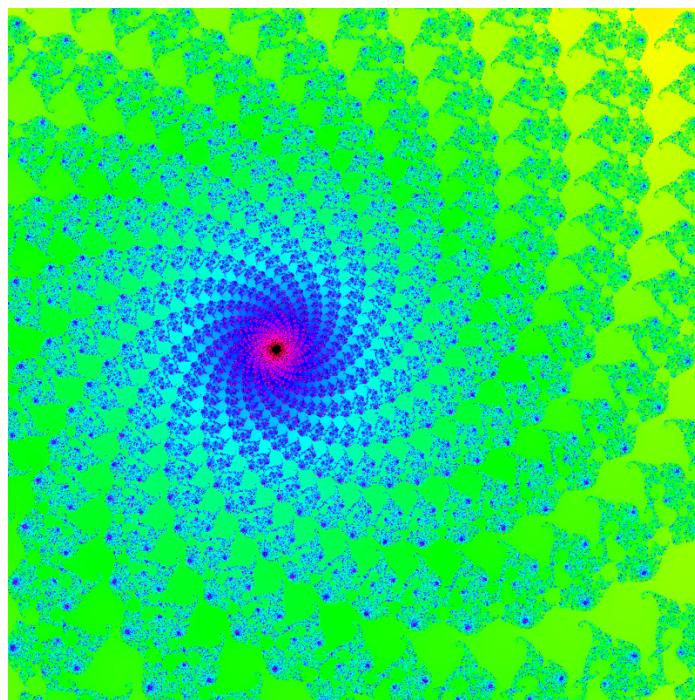


Figure 10 → `mandel -x 0.2869325 -y 0.0142905 -s .000001 -W 1024 -H 1024 -m 1000`

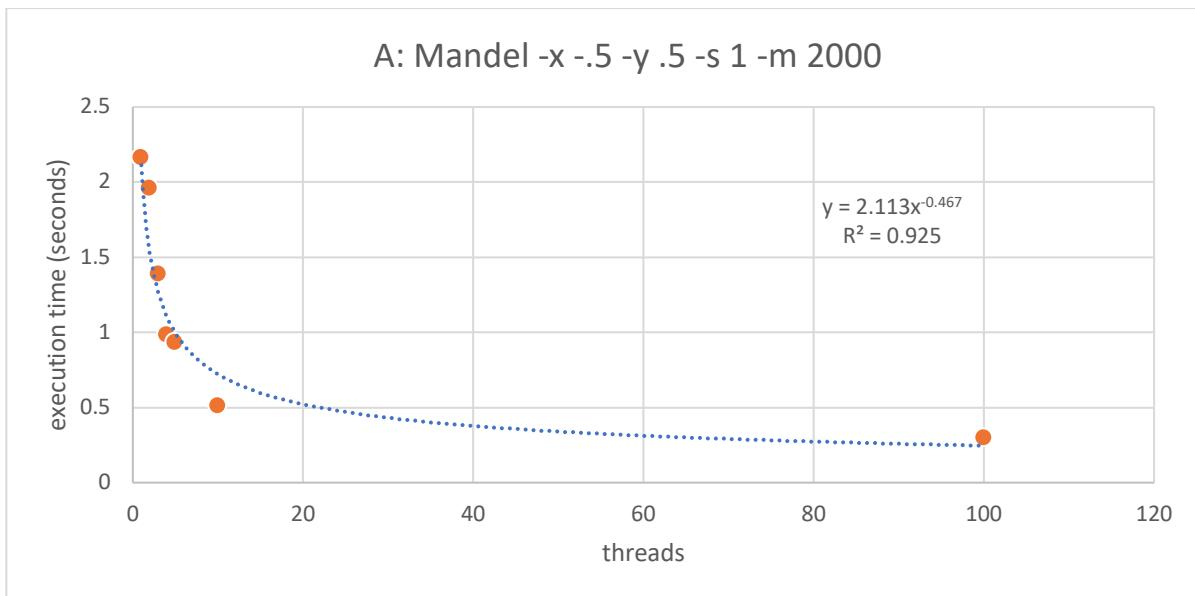
- A: mandel -x -.5 -y .5 -s 1 -m 2000

Figure 11 → ./mandel -x -.5 -y .5 -s 1 -m 2000

- A: mandel -x -.5 -y .5 -s 1 -m 2000

A: Mandel -x -.5 -y .5 -s 1 -m 2000	
threads	execution time (seconds)
1	2.161344
2	1.958613
3	1.388212
4	0.982542
5	0.931026
10	0.506999
100	0.293768

Table 2 → A: mandel -x -.5 -y .5 -s 1 -m 2000



- Curve Analysis:
 - The curve demonstrates an exponential decrease in execution time as the number of threads increases. Initially, the benefits from adding more threads are substantial, but as we continue to increase the thread count, the execution time's reduction rate decreases, hinting at a diminishing return.
- Optimal Number of Threads:
 - For Configuration A, the most significant improvement in execution time is observed when moving from 1 to 10 threads. Beyond 10 threads, the decrease in execution time is less pronounced, indicating that increasing the number of threads further might not offer substantial benefits. Thus, in this context, 10 threads can be considered as the optimal number.
- In conclusion, while multi-threading provides performance enhancements, it's essential to identify the optimal number of threads to achieve the best results without unnecessary overhead. For this specific task and based on the data, using 10 threads would be the most efficient choice.

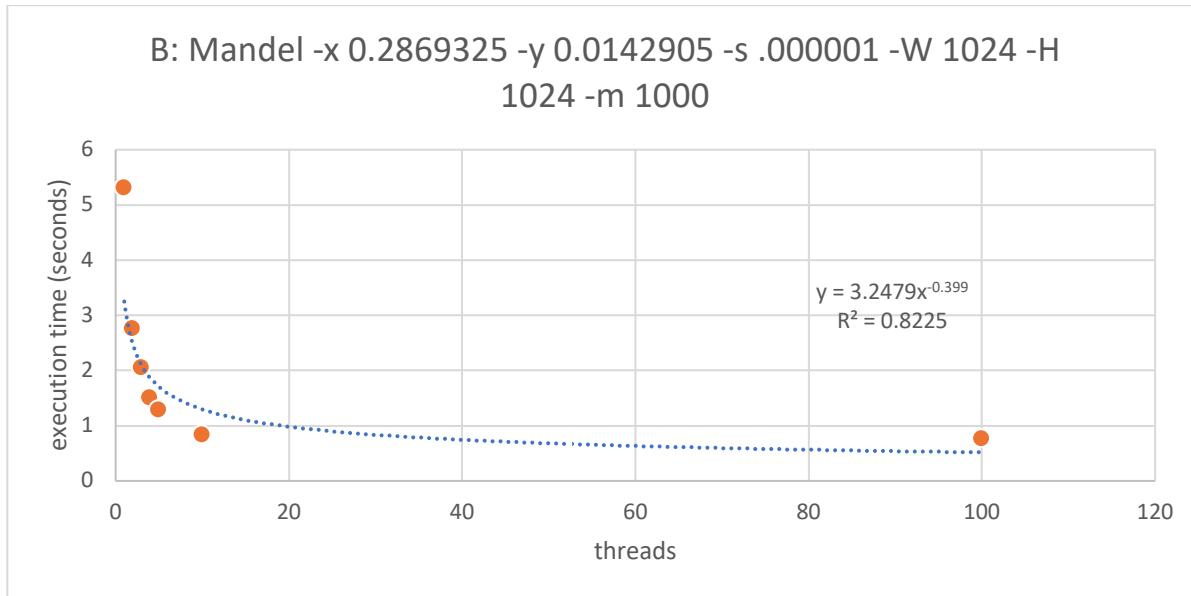
- B: mandel -x 0.2869325 -y 0.0142905 -s .000001 -W 1024 -H 1024 -m 1000

Figure 12 → ./mandel -x 0.2869325 -y 0.0142905 -s .000001 -W 1024 -H 1024 -m 1000

- B: mandel -x 0.2869325 -y 0.0142905 -s .000001 -W 1024 -H 1024 -m 1000

B: Mandel -x 0.2869325 -y 0.0142905 -s .000001 -W 1024 -H 1024 -m 1000	
threads	execution time (seconds)
1	5.306489
2	2.755431
3	2.03922
4	1.501214
5	1.282665
10	0.821996
100	0.756909

Table 3 → B: mandel -x 0.2869325 -y 0.0142905 -s .000001 -W 1024 -H 1024 -m 1000



- Curve Analysis:
 - The curve exhibits a power-law decrease in execution time as the number of threads increases. There's a pronounced drop in execution time when moving from 1 thread to a higher number of threads. However, as we further increase the thread count, the decrease in execution time becomes more gradual. This suggests that the benefits from additional threads diminish as the thread count increases.
- Optimal Number of Threads:
 - For Configuration B, the most significant reduction in execution time occurs between 1 and 10 threads. Beyond 10 threads, the improvement is less substantial, although there's still a noticeable decrease up to 100 threads. However, considering the diminishing returns and potential overhead associated with managing a large number of threads, using 10 threads seems to offer a balance between performance gain and computational efficiency.
- In summary, Configuration B shows that while increasing threads can reduce execution time, there's a point of diminishing returns. Based on the data presented, 10 threads appear to be the optimal choice for this task.

- Differences in Curve Shapes:
 - Curve A has a more pronounced decrease in execution time with increasing threads initially, followed by a plateau. Curve B, on the other hand, has a gentler decrease and starts plateauing earlier.
 - The difference in curve shape could be attributed to various factors:
 - Nature of the computation: The area of the Mandelbrot set being computed might have different complexities. The commands point to different regions of the Mandelbrot set to be computed, which can have varying levels of detail and complexity.
 - The size of the image to be computed: besides of the level of detail of the specific region of the Mandelbrot set that is computed for each command, there is also the size difference, where command B computes an image of 1024 by 1024 pixels whereas command A computes the default 500 by 500 pixels size. Here is a picture that compares the difference in size:

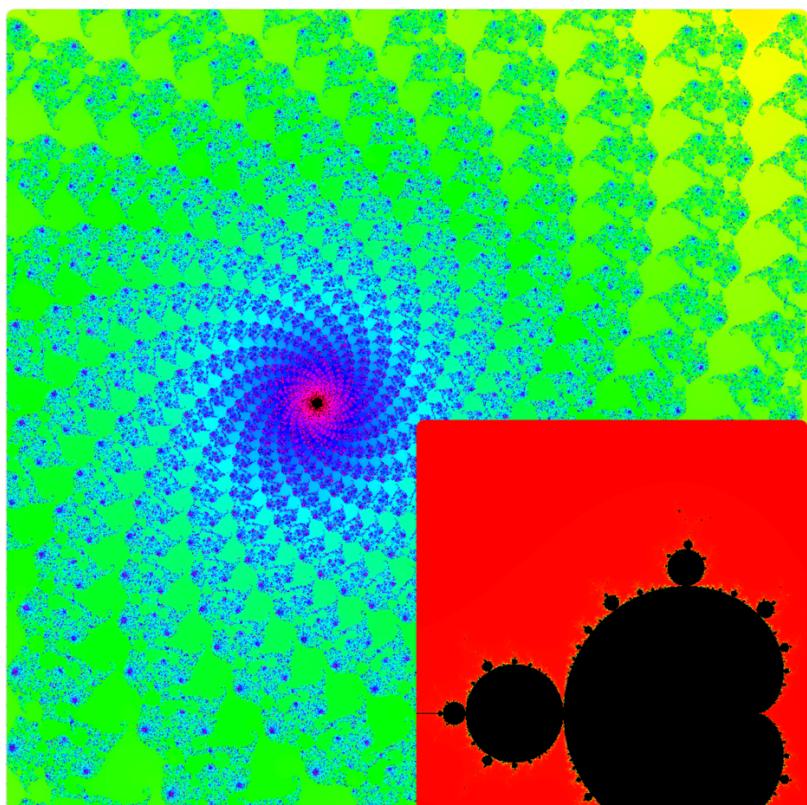


Figure 13 → 1024x1024 px vs. 500x500 px

- Conclusion:
 - The findings from the project accentuate the nuanced nature of parallel processing and multi-threading. Both techniques significantly expedite computation, with a notable decrease in execution time as processes and threads increase. However, the diminishing returns beyond certain points (5 processes or 10 threads) underscore a trade-off between performance gain and computational efficiency. While the optimal setup may vary depending on specific goals and machine configurations, a balanced approach, avoiding excessive parallelism, emerges as a prudent strategy for exploiting the benefits of parallel processing in generating Mandelbrot set images efficiently.